

# Pinpointing Resource Wastage with Profiling

Pengfei Su

Assistant Professor


Department of Computer Science and Engineering



# Resource Concerns are Everywhere



**Programs need to  
be efficient at all scales**



# Time is Money



Drop sales by 1%  
every 0.1s of latency



Drop traffic by 20%  
every 0.5s of latency



Drop visitors by 40%  
after 3s of latency

# Classic Profiling Techniques

- Identify hotspot — high resource utilization
  - ◆ Time
  - ◆ CPU usage
  - ◆ Cache misses
  - ◆ Memory allocations
- Limitations
  - ◆ Not tell if resources are “well spent”
  - ◆ Need significant manual efforts to investigate root causes

Pinpoint resource wastage instead of utilization

# Agenda

Profiling

CPU memory  
wastage

ICSE'19, FSE'19, ICSE'22,  
CGO'23, TACO'23, CGO'26

GPU memory  
wastage

ASPLOS'23, USENIX ATC'24,  
ASPLOS'26 (under review)

# Wasteful Memory Operations

Silent load

```
x = A[i];  
y = A[i];      y = x;
```

Silent store

```
A[i] = 10;  
x = A[i];  
A[i] = 10;
```

Dead store

```
A[i] = 0;  
A[i] = 10;
```

Two operations involved:  
one is **silent/dead**  
because of the **killing** one

## Solution: binary analysis

# A Motivating Example: SPEC CPU2006 Hmmer

- The compiler may do harm than good

Unoptimized

```
for (k = 1; k <= M; k++) {  
    mc[k] = mpp[k-1] + tpmm[k-1]; First store  
    if ((sc = ip[k-1] + tpim[k-1]) > mc[k])  
        mc[k] = sc; Second store
```

-O3 optimized

```
1 for (k = 1; k <= M; k++) {  
2     R1 = mpp[k-1] + tpmm[k-1];  
3     mc[k] = R1; Dead store  
4     if ((sc = ip[k-1] + tpim[k-1]) > R1)  
5         mc[k] = sc; Second store
```

The compiler: they “alias” to  
same memory location

# Necessity of Binary Analysis

```
1 for (k = 1; k <= M; k++) {  
2     R1 = mpp[k-1] + tpmm[k-1];  
3     mc[k] = R1;  
4     if ((sc = ip[k-1] + tpim[k-1]) > R1)  
5         mc[k] = sc;  
  
else  
    mc[k] = R1;
```

```
1 mov %r10,%rax,4), %ecx  
2 add 0x0(%r13,%rax,4),%ecx #mpp[k-1]+tpmm[k-1]  
3 mov %ecx, 0x4(%rdx) #assign mc[k]
```


Dead store

The value of `mc[k]` in memory location `0x4(%rdx)` is unused

```
10 mov %ecx, 0x4(%rdx) #assign mc[k]
```

Optimization: conditional check

# Hmmer's Speedup after Optimization



# User Inputs

- Rodinia-3.1 backprop

```
1 for (j = 1; j <= ndelta; j++) {  
2     new_dw = ETA * delta[j];  
3     w[k][j] += new_dw;  
4 }
```


Optimization

```
1 for (j = 1; j <= ndelta; j++) {  
2     if (delta[j] == 0) continue;  
3     new_dw = ETA * delta[j];  
4     w[k][j] += new_dw;  
5 }
```


Optimization: conditional check  
Speedup: 13%

# Type of Silent Loads

- Temporal silent load
  - ♦ Repeatedly load the same value from the same memory location





- Spatial silent load
  - ♦ Repeatedly load the same value from nearby memory locations




# Temporal Silent Load Detection

- Intercept every memory load to obtain current loaded value ( $V_{cur}$ )
- Employ shadow memory to save previous loaded value ( $V_{prev}$ )



# Spatial Silent Load Detection

- Intercept every memory load to obtain current loaded value
- Employ shadow memory to save previous loaded value
- Identify the memory range allocated for a data object
  - ◆ Static object: read the symbol table
  - ◆ Dynamic object: intercept malloc() family of functions and mmap()



# Case Studies

Program		SLoC	Inefficiency	% of silent load reduction	Speedup
Benchmark					
SPEC CPU2006	lbm	3K	Redundant computation	60%	1.25x
	imagick_r	274K	Redundant computation	51%	1.25x
	lavaMD	800	Redundant function calls	93%	1.39x
	srad_v1	600	Inefficient register usage	29%	1.11x
	srad_v2	200	Inefficient register usage	32%	1.12x
	particlefilter	600	Linear search	97%	9.8x
Application	Apache Avro-1.8.2	46K	Missing inline substitution	37%	1.19x
	Hoard-3.12	22K	Redundant computation	2%	1.14x
	MASNUM-2.2	121K	Linear search	36%	1.79x
	USQCD Chroma-3.43	929K	Missing inline substitution	7%	1.06x
	Shogun-6.0	546K	Missing inline substitution	2%	1.06x
	Facebook Stack RNN	2K	Redundant computation	15%	1.09x

# Agenda

Profiling

CPU memory  
wastage


ICSE'19, FSE'19, ICSE'22,  
CGO'23, TACO'23, CGO'26

GPU memory  
wastage


ASPLOS'23, USENIX ATC'24,  
ASPLOS'26 (under review)

# GPU Memory

- Problem 1: Fast-growing memory demand in DL
  - ◆ The trend of developing deeper and wider neural networks  
A trillion-parameter model with 3D parallelism requires 300+ A100 GPUs
- Problem 2: Slow-growing GPU memory budget
  - ◆ HBM must fit in small package space and chip periphery
  - ◆ Compared to increasing memory capacity, reducing memory latency is GPU vendors' top priority




Fast-growing memory demand



Slow-growing GPU memory budget


# GPU Memory

- Problem 3: Memory wastage is significant and pervasive across DL workloads




# Our Solution: DrGPUM

- A data-object-centric GPU profiler, including
  - ◆ Object-level analysis (e.g., arrays and tensors)
  - ◆ Intra-object analysis (e.g., individual elements in an array/tensor)




# Identified Patterns of Memory Wastage



# Late Deallocation

```
1. cudaMalloc(void*& 01, size)
2. cudaMemcpy(01, hostptr, size,
cudaMemcpyDefault)
3. cudaMalloc(void*& 02, size)
4. cudaMemcpy(02, hostptr, size,
cudaMemcpyDefault)
5. cudaMalloc(void*& 03, size)
6. cudaMemset(03, 0, size)
7. kernel1<<<block, grid>>>(01, 03)
8. kernel2<<<block, grid>>>(03)
9. kernel3<<<block, grid>>>(02, 03)
10. cudaMemcpy(hostptr, 03, size,
cudaMemcpyDefault)
11. cudaFree(01)
12. cudaFree(02)
13. cudaFree(03)
```



# Demo

- SimplyMultiCopy
  - ◆ A multi-stream program released with NVIDIA CUDA Toolkit

